

MarkLogic in the Cloud Workshop

Unit 3 – Pausing and Resuming Clusters


© 2018 MARKLOGIC CORPORATION

Goal: Pause and Resume a Cluster

- Pausing an AWS MarkLogic cluster
- Resuming the cluster again.

We will examine why and how to pause and resume an AWS MarkLogic cluster.

Auto Scaling Groups



SLIDE: 3

© 2018 MARKLOGIC CORPORATION

Amazon's Auto Scaling Group launches the desired number of MarkLogic server instances and ensures this number is maintained. Should a MarkLogic server's HealthCheck application server fail to respond, that instance is terminated and a new one replaces it.

The number of instances is a parameter in our CloudFormation templates. As long as the number of nodes is set to 1 or greater, there will be MarkLogic servers to respond to our requests. The Auto Scaling Group responds to our template parameters so scaling up or down is an easy process.

0 = hibernate

Number of Zones Total number of Availability Zones. 1 or 3.

Nodes per Zone Total number of nodes per Zone. Set to 0 to shutdown/hibernate

Availability Zone

The Availability Zones for VPC subnets and instances. Accept either 1 zone or 3 zones. In the order of Subnet 1, Subnet 2 and Subnet 3 (if applicable).

SLIDE: 4

© 2018 MARKLOGIC CORPORATION

A great benefit of Cloud Computing is that you only pay for what you use. If you are running a development server, or a site that doesn't need to be running 24/7, you can pause your entire cluster so you don't incur EC2 charges while it is not running.

The Auto Scaling Group and Managed Cluster features, along with CloudFormation, enable you to quickly pause your cluster. It is just as simple to resume the cluster and have your resources re-attach to all of your data, so that your cluster will be up and running where you left off.

Amazon's Auto Scaling uses the "Node per Zone" value to determine how many MarkLogic server instances are in each Availability Zone. A value of "0" essentially pauses the cluster since there will be no MarkLogic servers. Remember, your data and configuration is not discarded and will be available when new MarkLogic server instances start.



Unit 3 exercise – Pause & Resume

- Display the CloudFormation stack.
 - Update the CloudFormation template.
 - Set the number of instances to 0.
 - Check the status of the Cluster instances.
 - Update the CloudFormation template again.
 - Set the number of instances to 1.
- **DO**
 - Refresh to check the CloudFormation stack status.
 - Ensure each MarkLogic EC2 instance is available to Amazon's Elastic Load Balancer.

Questions?